

CAREER ASPIRATION

MSc Artificial Intelligence candidate with hands-on projects in GenAI evaluation, NLP extraction, cloud-deployed fintech analytics, and SQL/dbt data products. Seeking applied AI/ML roles building reliable tools for workflow automation, model evaluation, and data-driven decision support.

SKILLS

Programming	Python, SQL, Java, Bash
AI & ML	PyTorch, Hugging Face Transformers, scikit-learn, PySpark, feature engineering, model evaluation
GenAI & NLP	RAG evaluation, LLM evaluation, prompt engineering, structured extraction, red-team testing, guardrail checks
Data & Cloud	dbt, DuckDB, PostgreSQL, BigQuery, GCP, Cloud Run, Cloud Storage
Engineering	FastAPI, Streamlit, Pydantic, Docker, GitHub Actions, CI/CD, Git, monitoring
Analytics	A/B testing, CUPED, synthetic control, calibration, customer segmentation, dashboards
Languages	English (Fluent), Mandarin Chinese (Native), Japanese (JLPT N1)

SELECTED TECHNICAL PROJECTS

Customer Growth & Pricing Intelligence Platform <i>Python, SQL/dbt, DuckDB, BigQuery, Cloud Run, FastAPI, Streamlit, scikit-learn</i>	2026 <i>Fintech Data Science Project</i>
- Built synthetic fintech data product for activation, retention, CLV proxy, pricing scenarios, referral incrementality, experimentation, monitoring, and dashboards. - Modelled DuckDB/BigQuery dbt warehouse from Cloud Storage raw data to staging, intermediate, and mart layers, with 107 passing dbt checks. - Developed activation scoring workflow with calibrated models, explainability, segment checks, batch score extracts, model card, and monitoring. - Analysed onboarding and referral growth using SRM checks, CUPED, guardrails, difference-in-differences, synthetic control, and unit-economics framing. - Deployed public Streamlit dashboard, FastAPI contract, private Cloud Run API, scheduled scoring/monitoring jobs, CI, alerts, and operations runbook.	
Internal AI Agent Evaluation Lab <i>Python, FastAPI, Streamlit, Pydantic, pytest, Docker, GitHub Actions</i>	2026 <i>GenAI Evaluation/LLMOps Project</i>
- Built public synthetic evaluation lab for internal AI-agent reliability across RAG, structured extraction, safe refusal, tool governance, auditability, and monitoring. - Created reproducible benchmark with 24 runbook sections, 180 operations tickets, 296 golden cases, and 60 red-team cases. - Evaluated retrieval, citation coverage, issue/action accuracy, abstention accuracy, schema validity, red-team safe response rate, and tool block/audit rates. - Implemented Pydantic validation, trace IDs, audit events, monitoring snapshots, OTel-style span export, Streamlit dashboard, GitHub Pages reports, Docker, and CI. - Documented benchmark limitations, synthetic-data profile, safety boundaries, and follow-up work for stronger hand-authored cases and optional LLM extraction.	
LLM Event Extraction on GPU/HPC <i>Python, Qwen2.5-7B-Instruct, Hugging Face, NVIDIA A100, Slurm</i>	2026 <i>NLP/LLM Team Project</i>
- Contributed to team NLP project building a zero-shot LLM event extraction baseline on MAVEN and WikiEvents datasets. - Compared unconstrained and constrained-label prompting for trigger detection, event type prediction, and argument extraction. - Ran float16 inference with Qwen2.5-7B-Instruct on NVIDIA A100-SXM4-80GB using Hugging Face Transformers and Slurm. - Evaluated outputs using valid JSON rate, trigger accuracy, type accuracy, combined accuracy, and structured error analysis.	

EDUCATION

MSc Artificial Intelligence , University of Sheffield <i>Core modules: Scalable Machine Learning, Natural Language Processing, Parallel Computing with GPUs, Machine Learning, Data Science, Text Processing</i>	Sep 2025 — Sep 2026 Sheffield, UK
BSc Computer Science , Queen's University Belfast <i>Core modules: Data Structures and Algorithms, Software Engineering, Advanced Computer Architecture, Cloud Computing</i>	Sep 2021 — Jun 2024 Belfast, UK